Cyberguide: A Mobile Context-Aware Tour Guide

Gregory D. Abowd¹, Christopher G. Atkeson¹, Jason Hong¹, Sue Long^{1,2}, Rob Kooper¹ and Mike Pinkerton¹

¹Graphics, Visualization and Usability Centre College of Computing Georgia Institute of Technology Atlanta, GA 30332-0280

E-mail: {kooper,mpinkert,hong,abowd,cga}@cc.gatech.edu

² Wink Communications, Alameda, CA E-mail: sue.long@wink.com

Future computing environments will free the user from the constraints of the desktop. Applications for a mobile environment should take advantage of contextual information, such as position, to offer greater services to the user. In this paper, we present the Cyberguide project, in which we are building prototypes of a mobile context-aware tour guide. Knowledge of the user's current location, as well as a history of past locations, are used to provide more of the kind of services that we come to expect from a real tour guide. We describe the architecture and features of a variety of Cyberguide prototypes developed for indoor and outdoor use on a number of different hand-held platforms. We also discuss the general research issues that have emerged in our context-aware applications development in a mobile environment.

Keywords: Mobile computing, context-awareness, location-dependent applications, hand-held devices

1 Introduction

Future computing environments promise to free the user from the constraints of stationary desktop computing, yet relatively few researchers are investigating what applications maximally benefit from mobility. Current use of mobile technology shows a slow evolution from our current desktop paradigm of computing, but the history of interaction shows that the adoption of new technology usually brings about a radical revolution in the way humans use and view technology [11]. Whereas the effective use of mobile technology will give rise to an interaction paradigm shift, it is difficult to predict what that shift will be. We follow the advice of Alan Kay, therefore, and choose to predict the future by inventing it. Our approach is to think first about what activities could be best supported by mobile technology and then determine how the technology would have to work. This applications focus is important to distinguishing our work in mobile computing.

In April 1995, we formed the Future Computing Environments (FCE) Group within

the College of Computing and the Graphics, Visualization and Usability (GVU Center) at Georgia Tech to promote such an applications focus. Our group is committed to the rapid prototyping of applications that benefit from the use of emerging mobile and ubiquitous computing technologies. Quick development of these futuristic applications allows us to predict and shape what our everyday lives will be like when today's novel technology becomes commonplace.

Applications for a mobile environment should take advantage of contextual information, such as position, to offer greater services to the user. In this paper, we present the Cyberguide project, a series of prototypes of a mobile, hand-held context-aware tour guide. Initially, we are concerned with only a small part of the user's context, specifically location and orientation. Knowledge of the user's current location, as well as a history of past locations, are used to provide more of the kind of services that we come to expect from a real tour guide. We describe the architecture and features of a variety of Cyberguide prototypes developed for indoor and outdoor use on a number of different hand-held platforms. We also discuss the general research issues that have emerged in our experience of developing context-aware applications in a mobile environment. Some of these research issues overlap with those that we have considered in applying other applications of ubiquitous computing technology.

The general application domain which has driven the development of Cyberguide is tourism, but we have found it necessary to be even more focused in our research. The initial prototypes of Cyberguide, therefore, were designed to assist a very specific kind of tourist —a visitor in a tour of the GVU Center Lab during our monthly open houses. Visitors to a GVU open house are typically given a map of the various labs and an information packet describing all of the projects that are being demonstrated at various sites. Moving all of the paper-based information into a hand-held, position-aware unit provided a testbed for research questions on mobile, context-aware application development.

The long-term goal is an application that knows where the tourist is, what she is looking at, can predict and answer questions she might pose, and provide the ability to interact with other people and the environment. Our short-term goal was to prototype versions of Cyberguide on commercially available PDAs and pen-based PCs in which context-awareness simply meant the current physical position and orientation of the Cyberguide unit (and since it is hand-held, this locates the user as well). Position information improves the utility of a tour guide application. As the prototypes of Cyberguide evolve, we have been able to handle more of the user's context, such as where she and others have been, and we have increased the amount in which the tourist can interact and communicate with the place and people she is visiting.

1.1 Overview

This paper is an extended version of an earlier report on Cyberguide [7], we discuss the evolution of the Cyberguide design and prototype as well as what future research areas our experience has uncovered. We begin in Section 2 by describing scenarios for the use of context-aware mobile applications. In Section 3, we provide context for our research within the area of applications-centered mobile computing. The generic architecture of Cyberguide is explained in Section 4. We will describe in Section 5 the initial realization of the generic components of the Cyberguide architecture, a series of prototypes

developed for the Apple MessagePad. We will then describe in Section 6 how the initial indoor prototypes were extended for use outdoors and for greater interaction with the environment. We conclude in Sections 7 and 8 with a discussion of significant issues for context-aware applications development and how our past experience will influence our future development plans.

2 Scenarios for a mobile context-aware application

This section outlines some possible uses for future mobile context-aware applications. Some of these uses are currently being implemented and some are futuristic. We begin with our initial assumptions about what technology we expect Cyberguide to use. Tourists are usually quite happy to carry around a book that describes the location they are visiting, so a reasonable packaging would be in the form of a hand-held device. The ideal hand-held device will have a screen and pen/finger interface, access to substantial storage resources —possibly through an internal device such as a CD drive, or through substantial communication and networking resources (cell phone, pager, data radio interface) providing access to other storage servers (such as the Web)— an audio input and output interface with speech generation and potentially sophisticated voice recognition, and a video input and output interface. The video input (a video camera) could be pointed at the user to interpret user gestures, or pointed at the environment to interpret objects or symbols in the environment. The video output could be integrated into the main screen or be a separate video display device, such as an attached screen or heads up display on glasses worn by the user.

One major application of mobile context-aware devices are personal guides. Museums could provide these devices and allow users to take personalized tours seeing any exhibits desired in any order, in contrast to today's taped tours. In fact, many museums now provide portable devices for just such a purpose, but what we are envisioning is a device that would allow the tourist to go anywhere she pleases and be able to receive information about anywhere she is. Walking tours of cities or historical sites could be assisted by these electronic guidebooks. The hand-held devices could use position measurement systems such as indoor beacons or the Global Positioning System (GPS) to locate the user, and an electronic compass or inertial navigation system to find user orientation. Objects of interest could be marked with visual markers or active beacons or recognized using computer vision. Some objects, such as animals at a zoo or aquarium, might be difficult to mark but could be recognized with simple computer vision and some assistance from the environment (indications that this is the elephant cage, for example). The personal guide could also assist in route planning and providing directions. Some of these functions are currently being provided by automobile on-board navigation systems.

There are other ways to assist users. Consider a traveler in Japan that does not speak or read Japanese. The hand-held device could act as a pocket multilingual dictionary, actually speaking the appropriate phrase with the appropriate pronunciation to a taxi driver, for example (or even showing the appropriate Kanji and an associated map on the screen). A device that included video input or a scanner could assist in reading signs or menus. A device that could show stored images might be able to show a shopkeeper the desired object or favorite meal. Another more futuristic use is to assist the user by

recognizing faces at a cocktail party and reminding the user who people are.

Real-time communication allows a personal device to act as an agent for the user. A personal guide to a theme park could make reservations at particular rides, and alert the user when the reservation was available. The device could also tell the user which rides had the shortest lines. Similar approaches are currently being used for automobile traffic management in major cities.

An important application of context-aware devices is enhanced reality. A heads up display could provide "X-ray" vision for the user. While surveying a building for renovation, the location of hidden plumbing or electrical conduits could be indicated to the user, based on information from sensors and/or building plans. At an archeological site a visitor could be provided with various overlays indicating what used to be above the current ground level as well as what is below the current ground level.

Context-aware devices can also be used as tools. Simple sonar devices are used to make room measurements today. It would not take much to have a hand-held device that both videotaped and mapped a room along with user commentary. An ecological field study or an archeological dig could be assisted by a device that automatically recorded the context of a particular find, including noting the surrounding objects. Consider an electronic field guide that assisted the user in recognizing plants or insects.

One of the most interesting applications of context-aware devices is to support group interaction on a tour or in a classroom, for example. Participants in a live demonstration of some new technology could use their personal device to help steer the demo using majority voting or consensus among the viewers. Each participant could run a personalized version of the same demo by expressing their own choices. In this case context is which demo a participant is participating in or attending to, and the personal machine may switch to another context if it detects the user is attending to that context instead.

Many tourists take records of some sort of their travelling experiences, either by taking pictures or videos or by composing a travel diary. Imagine the possibilities if the recording of these experiences could be more efficiently and accurately recorded. A drive across the country could result in a trail superimposed upon a map, and clicking on the trail would reveal an image of what you could see at that moment —an automatically-generated spatial index into your memories.

These are but a few of the possibilities we can imagine that a context-aware application can provide for the tourist. We have investigated many of these possibilities already and report on them later.

3 Related Work

In thinking about and developing a location-aware application, we were greatly influenced by work such as the PARCTab at Xerox PARC [14], the InfoPad project at Berkeley [7], the Olivetti Active Badge system [14] and the Personal Shopping Assistant proposed at AT&T [3]. We wanted to build useful applications that might take advantage of the hardware developed in the PARCTab and InfoPad projects. We did not want to build our own hardware, so we have a different focus from all of these projects. There are a number of commercially available and relatively inexpensive hand-held units that would suffice

for our purposes, such as the Apple MessagePad with the Newton Operating System¹, a MagicCap² machine or a pen-based palmtop/tablet PC. We chose to work initially with the Apple MessagePad 100 with Newton 1.3 and pen-based PCs running Windows for Pen Computing 1.0. Each platform was available for \$ 150–500 with relatively powerful development environments. This low cost of hardware was critical to the success of Cyberguide because it made it possible to put a number of units in the hands of many students, all with different ideas that they were allowed to investigate.

For positioning, we considered the Active Badge system, but rejected it for reasons of cost and long-term objectives. The Active Badge system combines position detection with communication. For room-level granularity of position, this is reasonable since the communications range is on par with the position resolution. With Cyberguide, it is not clear that positioning and communication systems should always share physical resources. Certain versions of our prototype did; other prototypes did not. We provided for the separation of the wireless communications capabilities from the positioning system, so we could seek out more cost-effective solutions for both.

We tried to pay attention to the higher level conceptual design of Cyberguide, but we have not been as general in our handling of context-aware mobile objects as has Schilit [13].

4 Architecture of Cyberguide

From the beginning, we have viewed Cyberguide as a family of prototypes and not just a single prototype, so it made sense to think about a conceptual design, or architecture, that captured the essence of the mobile tour guide. We have divided the system into several independent components, or building, and have found it useful to present those components both in terms of the generic function and personified in terms of the people a tourist would like to have available while exploring unfamiliar territory. The overall system serves as a tour guide, but we can think of a tour guide as playing the role of cartographer, librarian, navigator and messenger. The services provided by these components are:

- Cartographer (Map Component) This person has intimate knowledge of the physical surroundings, such as the location of buildings, interesting sights within a building, or pathways that the tourist can access. This component is realized in our systems by a map (or maps) of the physical environments that the tourist is visiting.
- Librarian (Information Component) This person provides access to all of the information about sights that a tourist might encounter during their visit. This would include descriptions of buildings or other interesting sights and the identities of people associated with the areas. The librarian can answer specific question about certain sights ("Who works in that building?" or "What artist painted that

¹MessagePad and the Newton Operating System are registered trademarks of Apple Computer, Inc.

²MagicCap is a registered trademark of General Magic, Inc.

picture?" or "What other demonstrations are related to what I am looking at?"). This component is realized as a structured repository of information relating to objects and people of interest in the physical world.

- Navigator (Positioning Component) The interests of the tourist lie relatively close to their physical location. Therefore, it is important to know exactly where the tourist is, in order to show the immediate surroundings on the map or answer questions about those surroundings ("What am I looking at?"). The navigator is responsible for charting the location of the tourist within the physical surroundings. This component is realized by a positioning module that delivers accurate information on tourist location and orientation.
- Messenger (Communications Component) A tourist will want to send and receive information, and so the messenger provides a delivery service. For example, when visiting an exhibit or demonstration, the tourist might want to speak with the owner of the exhibit. If the owner is not present, the tourist can leave a message. In order to find out where other tourists are located, each tourist can communicate her current location to some central service that others can access. It might also be desirable to broadcast information to a set of tourists ("The bus will be leaving from the departure point in 15 minutes."). This component is realized as a set of (wireless) communications services.

The utility of this architectural decomposition for Cyberguide is that it provides an extensible and modular approach to system development. It is extensible because we can always add further services. For example, we have considered adding an historian whose purpose is to document where the tourist has been and what her reactions were to the things she saw. It is modular because it has allowed us to change the implementation of one component of the system with minimal impact on the rest of the system. For example, we have implemented different versions of the navigator and the librarian without having to alter the other components. Of course, these components are related in some ways; for instance, position information ultimately has to be translated into a location on the physical map. Defining standard interfaces between the components is the means by which we achieve separation between and coordination among the various components.

5 The Indoor Cyberguide

In this section, we describe how each of the separate modules in the conceptual architecture have been realized in the initial series of prototypes developed on the Apple MessagePad for use indoors during GVU open houses.

5.1 Map Component

The initial map module, shown on the left side of Figure 1, contains a map of the entire GVU Center. Passageways and demonstration stations (stars in Figure 1) are shown. Only a limited view of the lab can be seen at any given time. The user can scroll the map around and zoom in and out to see alternative views. There is an icon to show the user's




Figure 1: The map (left) and information (right) interfaces of the initial Cyberguide MessagePad prototype.

location on the map. Using information from the positioning module, we implemented automatic scrolling of the map. If desired, the user's position is updated automatically and the map is scrolled to ensure that the user's current position remains on the visible portion of the map.

5.2 Information Component

The information module (shown on the right side of Figure 1) contains information about each of the demos on display at the GVU open house. This includes abstracts of the project being demoed, background information on those involved with the project, as well as where to get further information. The location of each demo is marked on the map by a star. The user selects the star icon for a demo to reveal its name. Selecting the name brings up the information page for that demo. The user can also go directly into the information module and search for information for specific demo pages either by category or by project name.

One version of the information module was hard-coded, providing very fast response but requiring a recompilation every time demo information needed to be updated. Another implementation used Newton files, called soups, to store information. The use of soups avoided hard-coding data into the application and simplified demo information updates, but did not have adequate response time. Our third implementation of the information module used Newton Books, the Newton platform documentation viewer, to store the demo information. The use of Newton Books improved our access time consid-




Figure 2: Questionnaire using communications module for delivery.

erably, allowing for an automated information update process without requiring data be hard-coded directly into the application. Throughout all three versions of our information module, we were able to modify the information module independent of the development efforts of the other modules, validating the modularity of part of our design.

5.3 Communication Component

Our initial implementation of a communication module consisted of a wired Internet Appletalk connection from a Apple MessagePad through a Unix Appletalk Gateway. We designed an application level protocol on top of a public domain implementation of the Appletalk protocol for Solaris[4]. This allows us to open a connection-based Appletalk stream from the Apple MessagePad to a UNIX platform. We then invoked our gateway application to repacketize Appletalk packets into TCP/IP packets for transmission over the Internet. This allowed for TCP/IP connectivity from a Apple MessagePad via an Appletalk connection. We could then fetch HTML documents as well as send and receive e-mail. We utilized this functionality within Cyberguide by providing a questionnaire for users to complete, which was sent to the developers as an e-mail message. (see Figure 2)

5.4 Position Component

Position is the obvious starting point for a context-aware mobile device. We considered several methods for sensing the user location. Outdoors, continuous services, such as GPS, can be used. Indoors, however, GPS signals are weak or not available. We considered RF




Figure 3: IR positioning prototype (left) and the array of positioning beacons in the GVU Lab (right).

for indoor position measurement, but found off the shelf solutions too expensive.

One solution for an indoor positioning system was to use infrared (IR). Our first positioning system was based on using TV remote control units as active beacons, and using a special IR receiver tuned to the carrier frequency ($\approx 40\,\mathrm{kHz}$) of those beacons (Figure 3). A microcontroller (Motorola 68332) interfaced the IR receiver to the serial port of the Apple MessagePad. We deployed an array of remote controls hanging from the ceiling (Figure 3 right), each remote control acting as a position beacon by repeatedly beaming out a unique pattern. The 68332 translates the IR pattern into a unique cell identifier that is sent to the Apple MessagePad's serial port. As the tourist moves around the room and passes into the range of a new cell, the position (indicated by an arrowhead) is updated on the map. Keeping track of the last recorded cell location provides a good guess as to the location the tourist is heading, so we indicate an assumed orientation by pointing the position icon accordingly.

The remote control system is too expensive for large scale use as the cost of the 68332 microcontroller is roughly equivalent to that of the MessagePad.

6 Extending the initial prototype

The first Cyberguide prototypes were completed within 6 months. To test out the genericity of our architectural approach, we decided to develop further prototypes that altered






Figure 4: The outdoor Cyberguide (left) with GPS unit (right).

one or more of the major components described in Section 4 and increased overall functionality. We describe these extended prototypes here.

6.1 Outdoor positioning

There were several motivations for building a Cyberguide prototype for outdoor use (Figure 4). First, we wanted to use Cyberguide over a wider area than the relatively small GVU Center. We also wanted to test the modularity of our design by having to change critical features. The two features that were changed on this prototype were the underlying map and the physical positioning system. We obtained a different map and inserted that into the map module without any problems. For positioning, we replaced the IR positioning module with a Trimble GPS unit attached to the Apple MessagePad serial port. (see right side of Figure 4). The GPS unit sends a position in latitude and longitude which was then translated into a pixel coordinate representing the user's current position on the map.

The outdoor positioning system has been tested by two prototypes. We first built a proof of concept tour of the Georgia Tech campus (shown in Figure 4). We also developed a more functional outdoor prototype that covered three surrounding neighborhoods of the campus, described later

6.2 Alternate platforms

In order to verify the platform independence of our conceptual design, we initiated two separate efforts building pen-based PC versions of Cyberguide. These limited functionality PC versions were written using Borland's Delphi environment and Microsoft's Visual Basic. Both were initially installed on Dauphin DTR-1 palmtops running Pen for Windows Computing 1.0.

The Delphi version implements the map and information module. Web pages containing demo information are stored locally as database objects using a stand-alone Borland database engine. The information base is extracted from the collection of existing Web pages for GVU projects but stored locally. This information is viewed using a public domain Delphi HTML viewer. Though this provided a very fast response for information queries, it is a long-term disadvantage to have the information base stored locally. Too much in our environment is subject to change. A local and static database is only slightly more useful than a book. This approach to static information storage is used currently for on-board navigation systems on certain rental cars.

The Delphi prototype uses vector-based maps, allowing for arbitrary scaling and rotation of the map and well as path generation. While there are several sources for obtaining vector maps for outdoor regions, it is not so easy for indoors. Consequently, there is a trade-off between the easily generated but limited functionality of bitmap images and the highly functional but hard to generate vector maps for indoor use.

The Visual Basic prototype, shown in Figures 5–9 realizes all four components of Cyberguide, including two-way communications, which is discussed next. We implemented historical context by predicting when a user had visited a demo, based on time spent in the area of the demo and interaction with the map. In Figure 5, a visited demo is indicated on the map by a checkmark. There is also a separate panel that lists the demos visited. This information could be used, for example, to generate a summary of the day's visit to GVU open house and then mailed off to the visitor. We again made use of a publicly available HTML rendering component to display the project descriptions, still stored locally.

6.3 Increased communication

A number of interesting possibilities are enabled for the tourist in a wireless communication mode. We have spent a good deal of effort building an indoor, low-cost wireless communications infrastructure for Cyberguide. We have built a serial IR network using inexpensive modules from Sharp (the same modules used in the MessagePad). We have written UNIX server software and client software for both the MessagePad and PC. Figure 7 shows how our homemade network connects mobile units to the departmental network. We have defined a simple protocol to support three different kinds of messages: mailing out from a mobile unit to the network (as shown in Figure 8); broadcasting from the network to all mobile units (Figure 9); and updating positioning information. We implemented this protocol over serial IR instead of some other standard so that we could immediately use all platforms (UNIX, Newton and Windows). Given the appropriate hardware and protocol support (perhaps IRDA), we sencould provide the same functionality more robustly.




Figure 5: The main map interface of the PC Cyberguide. Checks on demo sites indicate the user has been to visit that demo already, indicating history-sensitive interface.




Figure 6: The information browser interface of the PC Cyberguide.



Figure 7: Home-made ir units allow cross-platform communication.



Figure 8: The mail interface of the PC Cyberguide. Addresses are automatically filled in and message templates are supplied upon request.




Figure 9: A demonstration of receiving a broadcast message from the Network to an individual Cyberguide unit.

6.4 Increased interaction with environment

All of our applications of Cyberguide so far have restricted the role of the tourist to browsing, but it is likely that as she visits some place, the tourist will want to keep a record of her experience as advice for herself or others later on. With this idea of increased interaction with the environment and recording in mind, we created another Cyberguide prototype, called CyBARguide, to assist a tourist in pursuit of refreshment at neighborhood establishments in Atlanta. This prototype covers approximately 12 square miles of midtown Atlanta, using multiple maps at varying levels of detail.

Figure 10 shows the map interface on the left and a view of a user-modifiable database for interesting establishments on the right. The tourist can indicate a desired destination and as she moves around, CyBARguide automatically chooses the map of the highest detail that contains both the traveler (indicated by a triangle in Figure 10), and the destination (the beer mug with the emboldened border in Figure 10). Along the way, if the tourist eyes another interesting establishment that is not currently highlighted on the map, it could be added. Each establishment has a user-modifiable database entry associated with it that reflects both objective (e.g., availability of parking, average price of drinks) and subjective (e.g., ambiance or other comments) information that can be used in the future to plan an evening's excursion, querying of a large amount of information and some minimal routing facilities. We also plan to make the data within the information module modifiable so the user can add personalized information including personal impressions that may be useful for future reference, a type of virtual graffiti. We envision the use




Figure 10: The CyBARguide interface. The left shows the interactive map indicating the user's location (the triangle) and the location of establishments previously visited (the beer mugs). The user modifiable database shown on the right supports the long-term development of touring information for a location.

of CyBARguide in a mode in which both tourist and proprietors are able to modify the information base. A simple searching capability can assist revellers in search of a certain kind of entertainment experience, and additional contextual information, such as the knowledge of where the traveller has been already, the time of day, and what special events (information provided by the proprietors) are currently scheduled, can be used to deliver suggestions for where to go.

7 Issues

Our experience over the last year developing versions of Cyberguide with different features on different platforms has given us a certain amount of insight into the important issues in developing mobile, context-aware applications. In this section, we summarize some of those issues.

Our prototyping efforts were iterative, modifying both hardware and software to improve functionality. Our primary focus, to assess the impact of mobile technology for a specific task, necessitated rapid and inexpensive prototyping. This motivated the use of inexpensive commercially available hardware. When choosing our hardware platform, we considered several mobile hand-held devices before deciding on the Apple MessagePad and

various palmtop PCs. One of the driving considerations in purchasing hardware was cost, because we wanted to provide as many units to individuals and groups to maximize the number of ideas investigated. The downside of this decision was a collection of applications that were not very robust, but that negative has been outweighed by the sheer number of options we have been able to investigate and prototype in little over a year.

Over the last year we have used Cyberguide for many GVU open houses. We gathered informal feedback through informal surveys, formal questionnaires and informal user comments. We also observed visitors as they tried to use Cyberguide to maneuver around the lab. During each iteration we incorporated the user feedback and our own reactions to what was good and what was bad into the next iteration on the design. While our major focus was to prototype a context-aware mobile application rapidly, we realize that little can be determined concerning the impact of such technology unless the technology is put in the hands of real users.

7.1 Coupling of Positioning and Communication

There is an interesting relationship between the positioning and communication systems. Systems such as PARCTab and the Active Badge rely on a close coupling of positioning and communication. This is because the location of any mobile entity is determined by the beacon which receives a communication from the entity. In Cyberguide, the indoor and outdoor positioning system worked by having the beacon inform the mobile entity where it was located. The disadvantage of this latter approach for Cyberguide is that only the mobile entity knows its location. For applications in which you want objects to know about the position of other objects, there must be some sort of communication. However, it can be impossible or undesirable to couple positioning and communication together. For example, if position is coming from GPS, then a separate means of communication must be used. In our current version of indoor IR positioning using the Sharp IR units, we can couple positioning and communication, but the range of the IR link is so limited (3 feet) that communication will be cumbersome. It makes sense to use a short range IR positioning system because position information can be localized to objects of interest. Communication, on the other hand, needs to be uniform throughout some space.

7.2 Communication Medium

We have been trying to implement communications services on commercial hand-held units. While communication is important to Cyberguide, there is no obvious appropriate choice for a wireless communication medium to suit our needs. It clearly is not a priority among the manufacturers of these units to provide high bandwidth, cost-effective wireless communications. There are many potential solutions for communications (IR, spread-spectrum RF, cellular packet, cellular modem), but the variety and quality of these services changes so much that manufacturers tend not to build communications devices into the units, for fear of premature obsolescence. Instead, they rely on third party communications solutions with standard interfaces (e.g., PCMCIA). In our experience, there is a need to have high bandwidth communication to the mobile unit and low bandwidth connection back to the network. We do not assume that the hand-held unit will always carry around with it the entire information base or map associated with the area the tourist is visiting. Rather, that information should be provided on demand and relative

to the position/orientation of the tourist.

7.3 Map Representation

We have experimented with bitmap and vector-based maps. The bitmap representation is easy to obtain (scanning) for any area and is relatively inexpensive to store and display. Scaling and rotation, however, are cumbersome with that representation. Since we were decorating the map to highlight places of interest that were not on the original bitmap, it was difficult to control the display of the decorations after scaling for a zoom in or out. Another problem with the bitmap representation is accuracy with respect to the real-world. In the outdoor version of Cyberguide, we noticed a drift on the positioning system for a certain region of the map attributed to the map itself being out of scale. Also, a bitmap representation is not suited to doing higher level map services, such as generating a path to direct the tourist to a location of interest.

A vector-based representation, on the other hand, is easier to handle in terms of manipulation and additional services such as way-finding, but was not a feasible solution on the interpreted platforms (Visual Basic and Apple MessagePad 100 with Newton 1.3 using the Newton Toolkit 1.5) because it was computationally overwhelming to manage the display (this may be solved with compiling capabilities of later versions of the Newton Toolkit). It is also more difficult to obtain a vector-based map for a large and detailed area. For example, when we built the Delphi prototype, it was not difficult to build a map tool to construct the vector-based map, but it would take a very long time to create a map of the GVU Lab with the detail we already had in the bitmap version. For outdoor use, there are already commercially available structured map databases for large areas, and these are being used in navigational systems in rental cars. However, the size of the map database prohibits local storage on hand-held units, so there is an even stronger argument for high bandwidth downstream wireless connectivity.

7.4 Cross-Platform Issues

We developed prototypes on multiple platforms to validate our claim of platform independence for the architecture. It was encouraging to see the same indoor positioning system work for both the Apple MessagePad and Visual Basic prototypes. Outdoor positioning is similarly cross-platform, relying on either a serial or PCMCIA interface.

For communication, we also see the need to standardize the interface and protocol. For information services, it is natural to want support for wireless TCP/IP to enable full Internet capability. The information browsing can then be treated as a Web browsing task, for example. There are efforts already to support TCP/IP for platforms such as the Newton platform, but the bandwidth does not yet support delivery of complex graphics, as we would need for map delivery. Without cross-platform wireless Internet support, we were forced to approximate the connectivity using simple wireless serial connections. These do not provide the reliability nor range that would be necessary for a commercial strength application, but they provided enough of the infrastructure to investigate functional capabilities for the user.

8 Future Work

We have worked for a year on Cyberguide. We have always tried to keep in mind the very long term goals of this kind of application, and many of those ideas have been expressed by others and by us in Section 2. The following is a list of features that we feel are most important for near-term future versions of mobile, context-aware applications such as Cyberguide.

8.1 Modifiable Information Base

As a tourist visits a place, she may read about prepared information, but she may also have her own thoughts and reactions to what she sees or may overhear someone else interpret an exhibit in an interesting way that she would like to record. Capturing relevant information along the way and adding that to the information base would be useful. We have been focusing on applications of Cyberguide that emphasize functionality based on knowledge of the user's physical context. Another major theme the FCE group has worked on involves automated capture of user experience to facilitate later access to a rich record of that experience. Most of our work in that area has involved the application of ubiquitous computing technology in the stationary environment of the university classroom [12, 1, 2] and others have applied it in the meeting room environment [9, 10]. We can see a hint of this capture/access problem appearing in CyBARguide as the traveller records impressions of an establishment that may influence their plans later on. It is a natural extension to add this kind of capture facility for the tourist.

We have experimented with capturing contextual information from a traveller to support the automatic authoring of a travel diary.³ As a user journeys around, a log of time and place is recorded. In addition, the traveller carries a digital camera or video and comments about a day's journey into a microphone. All of these streams of activity can be associated either by time or place, resulting in a rich multimedia record of a vacation. This record is also searchable, which supports queries such as, "What was the name of that charming gift shop we visited shortly after lunch in Albequerque yesterday?" With a suitable combination of positioning system, digital camera and wireless communication, we can create dynamic Web pages that catalog our travel experience for others to enjoy.

8.2 Increased Communications

Cyberguides wireless communications is unreliable and short-range, and that has prevented us from moving the information sources off of the hand-held unit. This is a high priority in our future work and we will most likely investigate third party wireless modems with TCP/IP support on the hand-held units. Since we are currently limited to what is in place from commercial vendors, we might now take the time to investigate research platforms such as the InfoPad.

³Take a look at http://www.gatech.edu/chow/roadtrip for a sample of what we have already done.

8.3 Improved Context Awareness

One way to view the capturing activity above is as a way to augment the system's understanding of the tourist's context by remembering what was interesting. We currently have a very limited notion of context in Cyberguide — physical location and crude orientation. We have experimented with capturing historical context (what sights have been visited already), but there are a number of other aspects of the tourist's context that can be useful. For instance, knowing where everyone else is located might suggest places of potential interest. Knowing the tourist's reaction to some exhibits would help in suggesting other related places of interest. Being aware of time of day or day of week may aid in more intelligent user queries of interesting attractions or activities. We feel that the secret to context-awareness is in doing it behind the scenes. The more that can be automatically captured and turned into context, the better. If the user has to explicitly inform the system about context information ("I am currently located here." or "This exhibit is boring to me." or "The museum is currently closed.") then the context is unlikely to be fully utilized.

8.4 Leveraging the Web

In the touring example, we found that much of the information we wanted to display was already available on the Web. Web browsing is now a natural information browsing metaphor. It makes sense to leverage this available information resource and mode of interaction for all of the information needs. The map module then provides hooks or links to the information on the Web, and it is delivered on demand, similar to when a user selects a URL on a browser. There have already been some research [3] and commercial attempts at providing mobile Web browsers.

8.5 Use of Vision

In the extreme case, we can think of a communion between the physical and electronic worlds, as suggested by work in augmented reality. Replace the hand-held unit with a pair of goggles and as the user wanders around, information about certain exhibits can be summoned up and overlaid on top of the actual image. Vision techniques can be used to augment the positioning system to inform the system and tourist what the tourist is looking at. We have experimented with vision systems as an extension to Cyberguide. Ultimately, we want to move toward personalized vision systems.

8.6 Ubiquitous Positioning System

Our current prototypes are exclusively indoor or outdoor, not both. This is mainly because we had no one positioning system that worked in both conditions. GPS is unreliable indoors and the IR-based beacon system is impractical for us to implement outdoors. We intend to integrate both positioning systems into one application, to allow a tourist to wander in and out of buildings and have Cyberguide automatically switch the positioning system.

8.7 Increased Multimedia Support

In the context of the GVU open houses, visitors often gather around a demo and stretch to watch the activity on a desktop machine as the researcher describes the related research. We would like to have Cyberguide units in the vicinity of the demonstration be able to pick up a live feed from the demonstration for display on their own unit. Initially, this could be achieved by a simple mirroring of the demonstration machine's display onto the hand-held unit. Ultimately, we envision the visitor "plugging in" to the demonstration and being able to control it from their hand-held unit.

9 Conclusions

We have described research we have been conducting over the past year prototyping mobile, context-aware applications. The main focus of our development effort has been to take an applications focus, tourism in this case, and determine what commercially available hardware can be used to deliver relevant information to the mobile user that is tuned to knowledge of their physical context. We described several iterations on the Cyberguide application, built to support tours through various venues. Our experience prototyping several variations of Cyberguide has clarified our own thoughts on how context-aware computing provides value to the emerging technology promising to release the user from the desktop paradigm of interaction. The experience also has helped to clarify some significant research issues in the continued development of mobile, context-aware applications for future computing environments.

Acknowledgments

The Cyberguide project is very much a group development effort. The authors would like to acknowledge the support of the several members of the Future Computing Environments Group at Georgia Tech, particularly Dietmar Aust, without whom this project would surely have failed. Gregory Abowd would like to acknowledge the numerous undergraduate students in his software engineering courses who put forth effort in defining and building many of the prototype systems. Of particular note is the nearly independent effort of David Chow, who was brave enough to experiment with the dynamic generation of a Web page to chart his trek across the country. Cyberguide was initially supported by a seed grant from Peter Freeman, Dean of the College of Computing at Georgia Tech. Subsequent effort on the project was support by BellSouth Intelliventures through the Industrial Affiliates Program in the GVU Center at Georgia Tech. Some of our equipment has been loaned to us from Apple Computer, Inc. and we would like to thank Jim Spohrer in particular for his continued support of our work.

References

[1] G. D. Abowd, C. Atkeson, A. Feinstein, Y. Goolamabbas, C. Hmelo, S. Register, and

- N. N. S. and Mikiya Tani. Classroom 2000enhancing classroom interaction and review. Technical Report GIT-96-21, Graphics, Visualization and Usability Center, Georgia Institute of Technology, September 1996.
- [2] G. D. Abowd, C. G. Atkeson, A. Feinstein, C. Hmelo, R. Kooper, S. Long, N. Sawhney, and M. Tan. Teaching and learning as multimedia authoring: The classroom 2000 project. In Proceedings of the ACM Conference on Multimedia Multimedia '96, 1996. To appear.
- [3] A. Asthana, M. Cravatts, and P. Krzyzanouski. An indoor wireless system for personalized shopping assistance. In L.-F. Cabrera and M. Sattyanarayanan, editors, Workshop on Mobile Computing Systems and Applications, pages 69-74. IEEE Computer Society Press, December 1994.
- [4] J. F. Bartlett. W4 the wireless World Wide Web. In L.-F. Cabrera and M. Sattyanarayanan, editors, Workshop on Mobile Computing Systems and Applications, pages 176– 178. IEEE Computer Society Press, December 1994.
- [5] B. Krupczak, M. Ammar, and K. Calvert. Multi-subsystem protocol architectures: Motivation and experience with an adaptor-based approach. In *Proceedings of IEEE INFOCOM'96*, 1996
- [6] S. Long, D. Aust, G. D. Abowd, and C. G. Atkeson. Rapid prototyping of mobile context-aware applications: The cyberguide case study. In Proceedings of the 1995 conference on Human Factors in Computing Systems CHI'96, 1996. Short paper.
- [7] S. Long, R. Kooper, G. D. Abowd, and C. G. Atkeson. Rapid prototyping of mobile context-aware applications: The cyberguide case study. In *Proceedings of the 2nd Annual International Conference on Mobile Computing and Networking*, November 1996. To appear.
- [8] A. C. Long, Jr., S. Narayanaswamy, A. Burstein, R. Han, K. Lutz, B. Richards, S. Sheng, R. W. Brodersen, and J. Rabaey. A prototype user interface for a mobile multimedia terminal. In *Proceedings of the 1995 conference on Human Factors in Computing Systems* CHI'95, 1995. Interactive experience demonstration.
- [9] S. Minneman, S. Harrison, B. Janseen, G. Kurtenbach, T. Moran, I. Smith, and B. van Melle. A confederation of tools for capturing and accessing collaborative activity. In *Proceedings of the ACM Conference on Multimedia — Multimedia* '95, November 1995.
- [10] T. Moran, P. Chiu, S. Harrison, G. Kurtenbach, S. Minneman, and W. van Melle. Evolutionary engagement in an ongoing collaborative work process: A case study. In *Proceedings of ACM CSCW'96 Conference*, 1996.
- [11] H. Rheingold. Tools for Thought. Prentice Hall, Englewood Cliffs, New Jersey, 1985.
- [12] N. N. Sawhney, G. D. Abowd, and C. Atkeson. Can electronic notebooks enhance the classroom? Technical Report GIT-96-06, Graphics, Visualization and Usability Center, Georgia Institute of Technology, February 1996.
- [13] W. N. Schilit. System architecture for context-aware mobile computing. PhD thesis, Columbia University, 1995.
- [14] R. Want, A. Hopper, V. Falcao, and J. Gibbons. The active badge location system. ACM Transactions on Information Systems, 10(1):91-102, January 1992.
- [15] R. Want, B. Schilit, N. Adams, R. Gold, K. Petersen, J. Ellis, D. Goldberg, and M. Weiser. The PARCTAB ubiquitous computing experiment. Technical Report CSL-95-1, Xerox Palo Alto Research Center, March 1995.